

# What is “mode”?

```
convolve2d(A, np.fliplr(np.flipud(w)), mode='valid')
```

valid | same | full  
~~same~~



↓ 1/p  
 $N = 5$   
 $K = 3$   
 $N - K + 1$   
 $= 5 - 3 + 1$   
 $= 3$   
 o/p  
 shrunk  
 in  
 size

- The movement of the filter is bounded by the edges of the image. The output is therefore always smaller than the input. **Input Image**

Filter?

0	1	2
0	1	2
0	1	2


3 <sub>0</sub>	3 <sub>1</sub>	2 <sub>2</sub>	1	0
0 <sub>2</sub>	0 <sub>2</sub>	1 <sub>0</sub>	3	1
3 <sub>0</sub>	1 <sub>1</sub>	2 <sub>2</sub>	2	3
2	0	0	2	2
2	0	0	0	1

Output Image

12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0

# Padding ( “same” mode)

- What if we want the output to be of the **same size** as the input image?
- Then we can use padding i.e. add imaginary zeros around the input.



$$\text{Input} = N = 5 (\rightarrow 1)$$

$$\text{Filter} = K = 3$$

$$O/P = 5$$

$$N - K + 1$$

$$= 5 - 3 + 1$$

$$= \cancel{4} + 1$$

$$\rightarrow = 5$$

# Even more padding! ( “full “ mode)

- We could extend the filter further and still get non-zero outputs.
- This is not very common these days.
- “full” padding
  - i. Input Length = N
  - ii. Kernel length = K
  - iii. Output length =  $N+K-1$

5

3

~~5+3-1~~

5+3-1

= 8-1

= 7

# Summary of Modes

- Input length =  $N$
- Kernel length =  $K$

It's here...



It's finally here!

24 ①

2

3

Mode	Output Size	Usage
Valid	$N - K + 1$	Typical
Same	$N$	Typical
Full	$N + K - 1$	Atypical

# A new perspective on Convolution

<https://i.insider.com/5e6f8386235c180f1533f1c2?width=800&format=jpeg&auto=webp>



Sliding Pattern Finder that passes through the entire image.



# How to view Convolution as Matrix Multiplication?

- Lets see 1D Convolution first because its easy to understand and 2D Convolution is just going to be a generalization of this.

Input image:  $a = [a_1, a_2, a_3, a_4]$

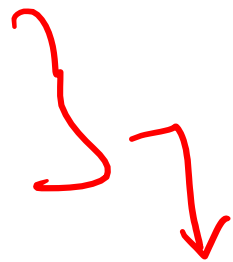
Filter:  $w = [w_1, w_2]$

Output image:  $b = a * w = [a_1w_1 + a_2w_2, a_2w_1 + a_3w_2, a_3w_1 + a_4w_2]$

$$a = (a_1, a_2, a_3, a_4)$$

$$w = (w_1, w_2)$$

CONVOLUTION



$$b = a * w = (a_1w_1 + a_2w_2, a_2w_1 + a_3w_2, a_3w_1 + a_4w_2)$$

# 1D Convolution in general

$$b_i = \sum_{i'=1}^K a_{i+i'} w_{i'}$$

- **Note:** It is very same as 2D Convolution's equation, just without 2<sup>nd</sup> index

# Matrix Multiplication


1D Array

- Can we do the same operation using Matrix Multiplication.
- YES!!!
- HOW?????????

Matrix Multiplication

Kernel \* 1/p

$$\begin{matrix} \begin{pmatrix} a_1w_1 + a_2w_2 \\ a_2w_1 + a_3w_2 \\ a_3w_1 + a_4w_2 \end{pmatrix} & = & \begin{pmatrix} w_1 & w_2 & 0 & 0 \\ 0 & w_1 & w_2 & 0 \\ 0 & 0 & w_1 & w_2 \end{pmatrix} & \begin{pmatrix} a_1 \\ a_2 \\ a_3 \\ a_4 \end{pmatrix} \\ \text{b} & & \text{w} & \text{a} \end{matrix}$$





# So what do we understand?

- By repeating the same filter again and again, we can do convolution without actually doing convolution.

$$\begin{pmatrix} a_1w_1 + a_2w_2 \\ a_2w_1 + a_3w_2 \\ a_3w_1 + a_4w_2 \end{pmatrix} = \begin{pmatrix} \underline{w_1} & \underline{w_2} & \underline{0} & \underline{0} \\ \underline{0} & w_1 & w_2 & \underline{0} \\ \underline{0} & \underline{0} & \underline{w_1} & \underline{w_2} \end{pmatrix} \begin{pmatrix} a_1 \\ a_2 \\ a_3 \\ a_4 \end{pmatrix}$$



# Problem with Matrix Multiplication?



# Problem with Matrix Multiplication?

① filter gets repeated multiple times

② It takes up a lot more space

③ Filter  $w$  was vector of size 2.

(In Matrix mult,  $w$  becomes a 2D Matrix of size 3x4.)

$$\begin{pmatrix} a_1w_1 + a_2w_2 \\ a_2w_1 + a_3w_2 \\ a_3w_1 + a_4w_2 \end{pmatrix} = \begin{pmatrix} w_1 & w_2 & 0 & 0 \\ 0 & w_1 & w_2 & 0 \\ 0 & 0 & w_1 & w_2 \end{pmatrix} \begin{pmatrix} a_1 \\ a_2 \\ a_3 \\ a_4 \end{pmatrix}$$

# Lets replace Matrix Multiplication with Convolution

MATRIX Mult. is BIG & SLOW.

∴ We will go for Convolution which is SMALL & FAST.



# Parameter Sharing / Weight Sharing

$$a = w^T \cdot x$$

whr  $a = o/p$  Activation

$x = i/p$  feature vector

$w$  = weight matrix.

Why what if I use the same weights ( $w_1, w_2$ ) again & again?

↳ ADV:

① Less parameters

② use lesser RAM

③ Computation becomes far more efficient.

∴ Convolution ~~takes~~ saves both SPACE & TIME by using less weights.



# Why do this?

Consider a fully connected ANN

EX: MNIST dataset

$28 \times 28 = 784$ -sized input vector

what if the image size is slightly larger & it is a color image?

EX: CIFAR-10 :  $32 \times 32 \times 3 = 3072$ -sized input vector.

Modern CNN Such VGG :  $224 \times 224$

If we are using full weight matrix, the # of features = 1,50,528 features

Modern HD image :  $1280 \times 720 = 2.8$  million features.

Too Large for a NN.




# Why do this?

Convolution  $\equiv$  Correlation.

FILTER  $\equiv$  Pattern Finder.

We want the same filter to look at all locations on the image.  
↳ This is called as Translational Invariance.



# Translational Invariance

- Suppose we are building a Dog vs cat classifier



Cat



Cat

If we use a fully connected NN, then the NN has to learn the weights for each of these positions separately.  
And yet after that, NN won't generalize that well bcz if we come across the same cat in a different position, NN shall fail to recognize it.



# Translational Invariance



Cat




Cat


∴ Having a Shared pattern finder makes more sense.

# Convolution on Color Images

Color! 3D :  $H \times W \times C$  ( $H \times \text{Width} \times \text{Color}$ )



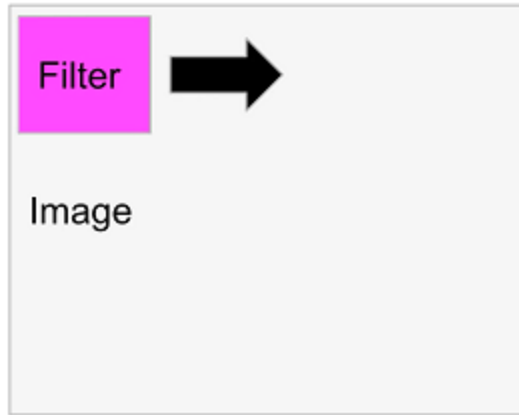
Original Color Image




Matlab RGB Matrix

# Convolution on Color Images


2-D



3-D



Note: same size in depth dimension



of p/image  
 $N - k + 1$   
 $5 - 2 + 1$   
 $= 4$

# Convolution on Color Images

2-D (2-D “dot product”) - a grayscale pattern-finder

$$(A * w)_{ij} = \sum_{i'=1}^K \sum_{j'=1}^K A(i + i', j + j') w(i', j')$$



E.g. if the filter is looking for red circles it won't match a green circle

3-D (3-D “dot product”) - a color pattern-finder


$$(A * w)_{ij} = \sum_{c=1}^3 \sum_{i'=1}^K \sum_{j'=1}^K A(i + i', j + j', c) w(i', j', c)$$

# What more?

3D  
~~~~

- Input Image:  $H \times W \times 3$ , which is a 3D vector
- Kernel:  $K \times K \times 3$ , which is a 3D vector
- Output Image:  $(H-K+1) \times (W-K+1)$ , which is a 2D vector
- We know that Neural Networks have a repeating structures (i.e. they have “uniformity”)
- Ex: For a Dense layer, if the input is 1D, then the output is also 1D and hence can be fed from one layer to another.
- But, in our case, we see that the output is 2D and input is 3D.
- Question: So how do we solve this?
- Answer: Lets see this now.



$I/P = 3D$   
 $O/P = 2D$



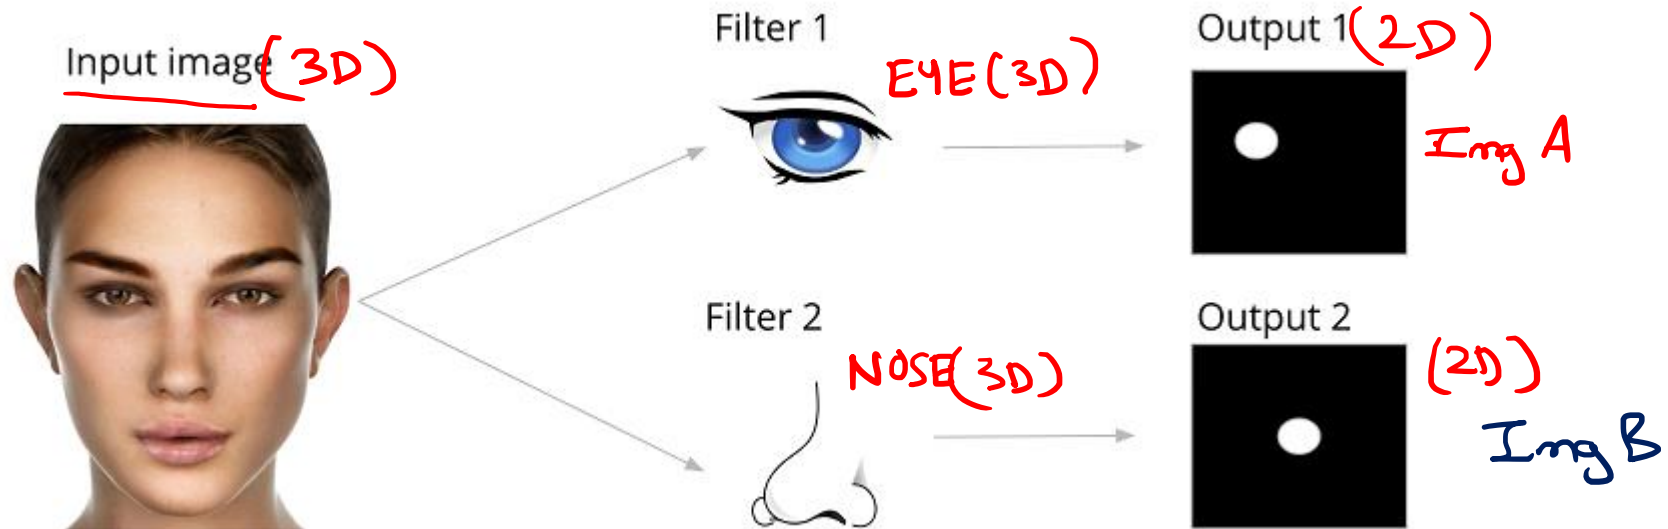
$I/P = N(2D)$   
~~~~  
Filter =  $k$   
 $O/P = N - k + 1$   
(2D)  
~~~~

$$(A * w)_{ij} = \sum_{c=1}^3 \sum_{i'=1}^K \sum_{j'=1}^K A(i+i', j+j', c) w(i', j', c)$$

# Multiple Features




eye, nose, lips,  
eyebrows




# Multiple Features

- Consider we have an image A with dimensions:  $H \times W \times 3$  <sup>3D</sup>
- If we use the “same” mode, and apply one filter, then we get output, let's say B1 with dimensions:  $H \times W$  <sup>2D</sup> eye
- If we use the “same” mode, and apply one more filter, then we get output, let's say B2 with dimensions:  $H \times W$  <sup>2D</sup> nose
- If we stack up B1 and B2 together, we get a new output image, let's say B with dimensions:  $H \times W \times 2$ , i.e. a 3D image as our original image image



# Multiple Features





# Convolution in Deep Neural Network

- Lets vectorize this operation, we don't need to do each color convolution separately

$$B = A \overset{\text{convolve}}{*} w$$

*A = Input Img, w = filter*

*B = Output Img*

$$\text{shape}(A) = H \times W \times C_1$$

$$\text{shape}(w) = C_1 \times K \times K \times C_2$$

$$\text{shape}(B) = H \times W \times C_2$$

$$B(i, j, c) = \sum_{i'=1}^K \sum_{j'=1}^K \sum_{c'=1}^{C_1} A(i+i', j+j', c') w(c', i', j', c)$$



# So What is this 3<sup>rd</sup> dimension in the Output image?

3<sup>rd</sup> is certainly not the color.

Input to NN:  $H \times W \times 3$

After subsequent convolutions, we get  $H \times W \times (\text{arbitrary}\#)$

3<sup>rd</sup> dimension = feature maps.



# Convolution Layer

$$\begin{aligned} \textit{Conv layer} : & \quad \underline{\sigma}(W \overset{\text{convolution}}{*} x + \underline{b}) \\ \textit{Dense layer} : & \quad \underline{\sigma}(W^T \cdot x + \underline{b}) \end{aligned}$$

# Shape of the bias

- In a Dense layer, if  $W^T x$  is a vector of size  $M$ ,  $b$  is also a vector of size  $M$
- In a Conv layer,  $b$  **does not** have the same shape as  $W * x$  (a 3-D image)
- Technically, this is not allowed by the rules of matrix arithmetic
- But the rules of broadcasting (in Numpy code) allow it
- If  $W * x$  has the shape  $H \times W \times C_2$ , then  $b$  is a vector of size  $C_2$ 
  - One scalar per feature map

$$\text{Conv layer : } \sigma(W * x + b)$$

$$\text{Dense layer : } \sigma(W^T x + b)$$

# How much do we save? CONVOLUTION

Has convolution saved my time & space?

Input img:  $32 \times 32 \times 3$

Filter:  $3 \times \underline{5 \times 5} \times 64$  (i.e. 64 feature maps)

Output img:  $28 \times 28 \times 64$  (assuming mode = 'valid')

why  $\downarrow$  28?  $N - k + 1$

$$H - k + 1 = 32 - 5 + 1 = 28$$

$$\therefore \# \text{ of parameters (ignore the bias)} = 3 \times 5 \times 5 \times 64 = \underline{\underline{4800}}$$



# How much do we save? MATRIX MULTIPLICATION

Flattened Input img =  $32 \times 32 \times 3 = 3072$

—||— O/p img =  $28 \times 28 \times 64 = 50176$


Weights Matrix =  $3072 \times 50176 = 1,54,140,672 \approx 154$  million

Compared to Convolution, we see that it takes  $\sim 32000$  times more parameters.

∴ Huge RAM.  $\rightarrow$  Suboptimal


# How are Convolution filters found?

- Since convolution is just a part of some neural network layer, it's easy to conceive of how the filters will be found
- Initially, we looked at convolution as an image *modifier* (blur, edge)
- Now, we see it as a pattern finder / shared-parameter matrix multiplication / feature transformer
- In other words,  $W$  will be found the same as before, automatically!
- Still gradient descent, `model.fit()`



*Conv layer* :  $\sigma(W * x + b)$

*Dense layer* :  $\sigma(W^T x + b)$






# CNN Architecture

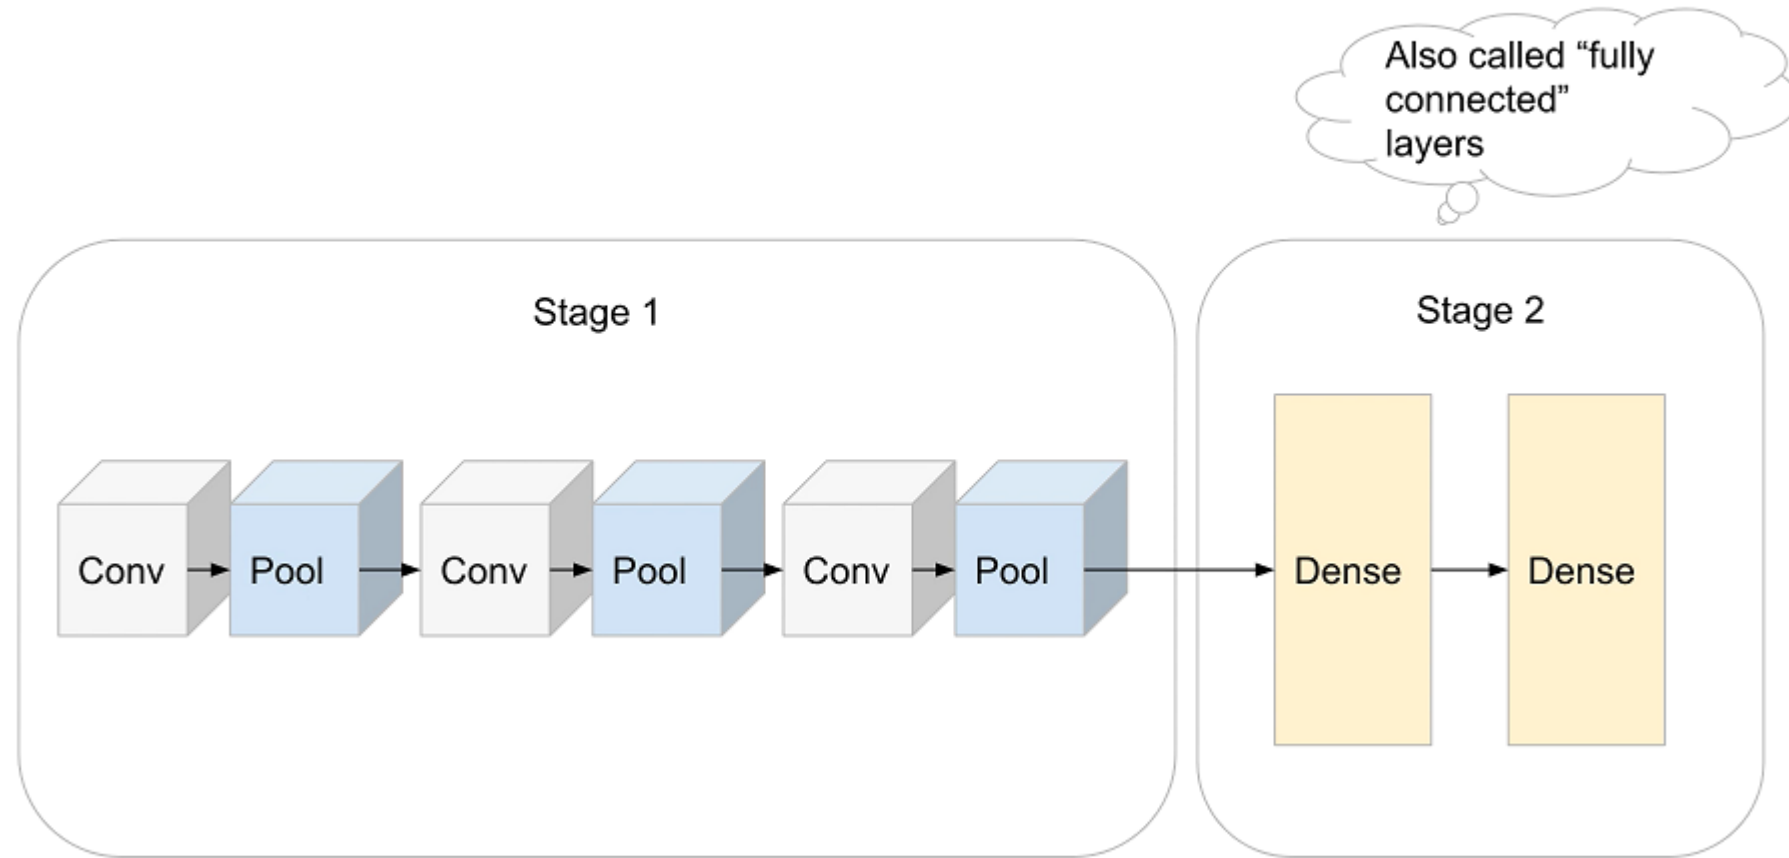


# CNN Architecture

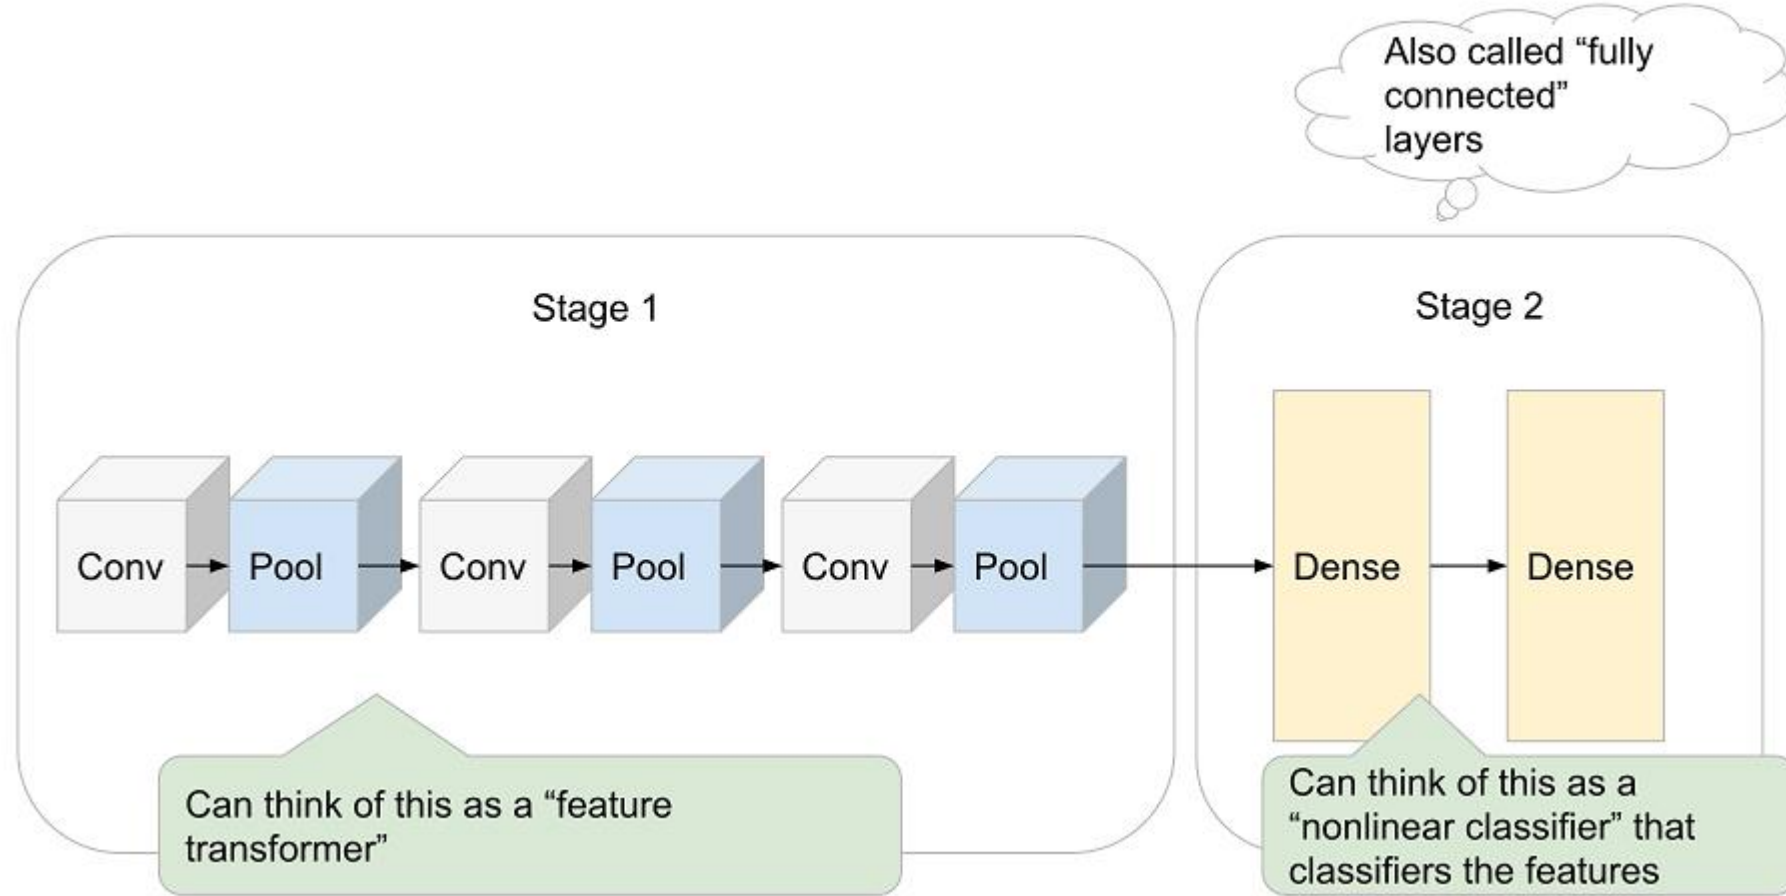
- Now that you understand how exactly a convolution layer works including the bias term and activation function we can now consider the architecture of a convolutional neural network and why it's that way.
- So as a little bit of a history lesson, modern CNN is essentially all originated from the same model, the LeNet.
- This is named after Yann LeCun, one of the original Deep Learning pioneers, along with Geoff Hinton and Yoshua Bengio.



# Typical CNN



# Typical CNN




**Pooling** means Downsampling  
i.e. output a smaller image from a bigger one.

eg: If i/p image =  $100 \times 100$   
then if pool-size = 2  
then my o/p image =  $50 \times 50$   
∴ Downsample by 2.




POOLING




# Types of Pooling

- There are two types of Pooling:
  1. Max pooling
  2. Average pooling
- Which one to use is a Hyperparameter choice.

# Max Pooling



# Average Pooling



# Why to use Pooling?

- Practical: If we shrink the image, we have less data to process.
- Translational Invariance: I don't care where in the image the feature occurred, I just care that it did.

