

Refer NB “6 CNN\_CIFAR.ipynb”

Results were not too good. OK OK....

# Data Augmentation

Improve our model  $\Rightarrow$  Data Augmentation

Translation Invariance : Humans can immediately identify this, but it is difficult for a computer / NN.

Cat



Same Cat




# Why data is important?

ML  $\Rightarrow$  Tabular Data

X					Y
-	-	-	-	-	-

500x6  
Grant: Yes/No  
email: Spam/Ham  
} you cannot create ur own data

If I try to create the new data, it will be contaminated with my own biased opinions.



# Why data is important?

Its perfectly fine to invent new data with images. Its meaning still remain the same.  
And no biases will added.




Cat



Cat (Tr-Invar.)

Cat (Rotation)



# Problem with this approach

## Data takes up space

The more data I invent, the more space it takes up!


There are an endless number of ways I can invent new data

Shift to the left by 1, 2, 3, 4, 5, ... pixels

Right / up / down also

Rotation

*Hyperparameter Optimization!!!*



# This can be done all automatically!

Have you ever considered how to rotate an image?


Use Tensorflow's Keras API instead

We need to know about generators / iterators

Data Augmentation  
2hrs 54min YouTube

Gen/Iterators: Play/Pause Button (A Song)

Walk (earphones - songs)  $\xrightarrow{\text{Play}}$  Uncle  $\xrightarrow{\text{Pause}}$   $\xrightarrow{\text{Play}}$  Friend  $\xrightarrow{\text{Pause}}$  ... repeat



# Generators / Iterators

Loop from 0...10: `for i in range(10)`

In Python 2, `range(10) = [0,1,2,3,4,5,6,7,8,9]` # a list!

In Python 2, use `xrange(10)` # does NOT create a list

In Python 3, `print(range(10))` yields `range(0,10)`

◦ Not a list!

`for i in range(10):`  
`print(i)`

# Create your own generator

- Can you write your own function to do something like this?

```
for x in my_random_generator():  
    print(x)
```

```
# 0.06654313  
# -0.68315371  
# 1.46795401  
# -0.9017639  
# 0.77572637
```

for x in [1, 2, 3, 4, 5]:  
 print(x)

o/p:  
1  
2  
3  
4  
5

Gun (6 bullets)

yield x  
1  
yield x  
2  
yield x  
3  
yield x  
4  
yield x  
5



# Yield

```
def my_random_generator():  
    for i in range(10):  
        x = np.random.randn()  
        yield x
```

Notice: no list is ever created

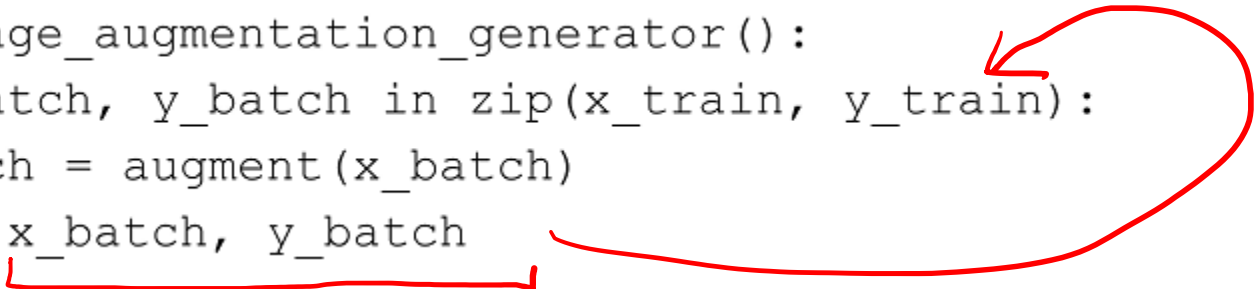
All values do not need to be stored in memory simultaneously

# Apply this to Data Augmentation

You could similarly generate augmented data on the fly

Conceptually, think of it like this

```
def my_image_augmentation_generator():  
    for x_batch, y_batch in zip(x_train, y_train):  
        x_batch = augment(x_batch)  
        yield x_batch, y_batch
```



# How does it work in tf.keras?

```
from tensorflow.keras.preprocessing.image import  
ImageDataGenerator
```

```
data_generator = ImageDataGenerator(  
    width_shift_range=0.1,  
    height_shift_range=0.1,  
    horizontal_flip=True)
```




Other args: rotation\_range, width\_shift\_range,  
height\_shift\_range, brightness\_range, shear\_range,  
zoom\_range, horizontal\_flip, and vertical\_flip

# How does it work in tf.keras?

```
data_generator = ImageDataGenerator(...) Prev. slide  
train_generator = data_generator.flow(  
    x_train, y_train, batch_size)
```

```
steps_per_epoch = x_train.shape[0] // batch_size  
r = model.fit_generator(  
    train_generator,  
    steps_per_epoch=steps_per_epoch,  
    epochs=50)
```




Car hor. flip = True  
vert. flip = 2

# Business as usual

`fit_generator` returns history, so plot loss per iteration, etc.

```
# Plot loss per iteration
import matplotlib.pyplot as plt
plt.plot(r.history['loss'], label='loss')
plt.plot(r.history['val_loss'], label='val_loss')
plt.legend()
```


<matplotlib.legend.Legend at 0x7f92f7ca6908>



# Batch Normalization

$x_{\text{train}} / 255.0$   
 $y_{\text{train}} / 255.0$

- Early on, we noted that it's important to normalize/standardize data before passing it into algorithms like linear/logistic regression → Standard/MinMax/Robust Scaling
- Problem: because this operation is done only on input data, only the first layer sees normalized data (after being transformed by Dense layer it's no longer normalized)



$$z = (x - \mu) / \sigma$$

∴ Go for Batch Normalization.


# Batch Normalization

- To start, recall: In Tensorflow, when we call `model.fit()`, we are doing **batch** gradient descent

```
① for epoch in range(epochs): // looping through each epoch
  ② for x_batch, y_batch in next_batch(x_train, y_train): // we look at a chunk of data
    w ← w - learning_rate * grad(x_batch, y_batch) on each step & do G.D.
                                                    wrt. this chunk/batch of
                                                    data.
```

# Batch Normalization

- What if we had a *layer* that would look at each batch, calculate the mean and standard deviation on the fly, and standardize based on that?



$$z = (x - \underset{\text{mean}}{\mu}) / \underset{\text{S.D}}{\sigma}$$



# Batch Norm as Regularization

Problem of Overfitting ↓


Can help with overfitting

Since every batch is slightly different, you'll get a slightly different  $\mu_B, \sigma_B$

They are not the true mean / std of the whole dataset


This is essentially noise, and using noise during training makes the neural network impervious to noise

- Rather than fitting to the noise!  
(a.k.a. "overfitting")



# Where is Batch Norm used?

*In b/w Convolutional Layers*



Refer NB “7  
CNN\_CIFAR\_Improved.ipynb”



# CNN Completed!!!

# DEEP LEARNING – Recurrent Neural Network

Trainer: Dr. Darshan Ingle.


# Recurrent Neural Network

Repeating

O/P

HL

I/P




Same




# Recurrent Neural Network (RNN)

Classify what kind of event is happening at every point in the Movie.




RNN have loops.



I/p

# RNN



An unrolled recurrent neural network.




~~#~~ Text Processing ~~#~~

# Why RNN?

Sequence:

Stream of data which are independent.  
eg: Time Series, Pieces of strings,  
conversation, etc.




CNN don't perform well when the i/p data is interdependent in  
a sequential pattern.

# Example

Predict what is the next letter?

Word: namaste  
n a m a s t e  
          ↑  ↑          ↑  
           $t-1$   $t$            $t$



Formula: 
$$h_t = f(h_{t-1}, x_t)$$

whr  $h_t$  = new state

$h_{t-1}$  = previous state

$x_t$  = input at time  $t$ .

Let us assume that we

are using AF tanh

$$h_t = \tanh(W_{hh} \cdot h_{t-1} + W_{xh} \cdot x_t)$$

Final o/p: 
$$y_t = W_{hy} \cdot h_t$$